

CrocoLakeTools: A Python package to convert ocean observations to the parquet format

Enrico Milanese¹, David Nicholson¹, Gaël Forget², and Susan Wijffels¹

¹ Woods Hole Oceanographic Institution, United States of America ² Massachusetts Institute of Technology, United States of America

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [Open Journals](#)

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Investigations of the ocean state are possible thanks to the ever growing number of measurements performed with multiple instruments by different research missions. The vast and heterogeneous efforts have brought the community to define data storage conventions (e.g. CF-netCDF) and to assemble collections of datasets (e.g. the World Ocean Database). Yet, accessing these datasets often requires the usage of specialized tools, is generally inefficient over the cloud, and remains a major obstacle to new users in terms of pre-required knowledge and time. CrocoLakeTools is a Python package that addresses those challenges by providing workflows to convert various oceanographic datasets from their original format to a uniform parquet dataset with a shared schema.

Statement of need

CrocoLakeTools is a Python package to build workflows that convert ocean observations from their original formats (e.g. netCDF, CSV) to parquet. CrocoLakeTools takes advantage of Python's well-established and growing ecosystem of open tools: it uses dask's parallel computing capabilities to convert multiple files at once and to handle larger-than-memory data. dask is already well-integrated with xarray and pandas, two widely used Python libraries for the treatment of array and tabular data, respectively, and with pyarrow, the API to the Apache Arrow library which is used to generate the parquet dataset.

Parquet is a data storage format for big tidy data which presents several advantages: it is language agnostic (it can be accessed with Python, Matlab, Julia and web development technologies); it offers faster reading performances than other tabular formats such as CSV; it is optimized for cloud systems storage and operations; it is widespread in the data science community, leading to a multitude of freely accessible tools and educational material.

CrocoLakeTools was developed with the goal of building and serving CrocoLake, a regularly refreshed database of sparse in-situ oceanographic observations that are pre-filtered to contain only quality-controlled measurements. CrocoLakeTools was designed to be used by researchers, engineers and data scientists in oceanography and to be accessed by the wider oceanographic community.

State of the field

In-situ measurements of ocean characteristics are collected by a wide range of projects at different scales: from small research groups to international consortia, from time-limited regional explorations to continuous global observation arrays. This has led to a highly heterogeneous landscape of products that often differ for schema (variable names, data types), data format,

and access services. As a result, ingesting and analyzing data from different sources, e.g. for research purposes, presents technical and logistical challenges. The need for homogenized ocean data products has then prompted several efforts, in particular when it comes to sparse observations, which are characterized by discrete measurements that are generally irregular and sparse in time and space.

One of the most comprehensive projects to date is the World Ocean Database (WOD) (Mishonov et al., 2024), which contains more than 40 variables gathered from several sources (buoys, gliders, floats, bathythermographs, etc.). The earliest record dates back to 1772, the data is quality-controlled, and the database is regularly updated with the latest measurements available. The data is public and freely available through different interfaces (WODselect, THREDDS, HTTPS, FTP) in ASCII, Comma Separated Value (CSV) and netCDF formats. WOD is maintained by the National Centers for Environmental Information (NCEI) of the United States' National Oceanic and Atmospheric Administration (NOAA).

Copernicus Marine Service is an initiative of the European Union that provides a wide range of ocean data products: physical, biogeochemical and sea ice data at regional and global scales, obtained from numerical models, in-situ observations and satellite observations. Among Copernicus' products, CORA (Szekely et al., 2019) (Coriolis Ocean database for ReAnalysis) provides quality-controlled in-situ temperature and salinity measurements gathered from several global and regional networks.

The International Quality-controlled Ocean Database (IQuOD) (Team, 2018) is an effort by the oceanographic community "with the goal of producing the highest quality and complete single ocean profile repository along with (intelligent) metadata and assigned uncertainties". It includes subsurface ocean profiles of several variables, paying particular attention to temperature measurements. The database is prepared by NCEI, and it is freely available in netCDF format through multiple channels (THREDDS, HTTPS, FTP).

argopy (Maze & Balem, 2020) is a Python library that facilitates access and manipulation of data from the real-time observing system Argo (Wong et al., 2020). The user provides some filters (e.g. coordinate and time ranges, measurement name, etc) and argopy retrieves the relevant data from the Argo Global Data Assembly Center (GDAC). ArgoData.jl (Gael Forget, 2025) and OceanRobots.jl (Gael Forget & Gebbie, 2025) are Julia packages that provide similar functionalities to argopy.

Argovis (Tucker et al., 2020) is a REST API and web application hosted at the University of Colorado, Boulder (United States) and serving profile data in JSON format from the Argo program, and from ship and drifters missions.

oce (Kelley et al., 2022) is a package written in R providing multiple functions to access oceanographic observations stored in a variety of formats. It focuses on datasets for physical oceanography with the goal of providing more tools in the future. argoFloats (Kelley et al., 2021) is a derived package that focuses on fetching and manipulating Argo data.

The Ocean Data Platform by the non-profit HUB Ocean is among the youngest projects in the community, and it allows to find and access datasets from a catalog. The user can interact with the platform through different interfaces (SDK, REST API, OQS, JupyterHub workspaces), loading the datasets in tabular format.

The above efforts often serve data in ASCII, CSV, JSON, or netCDF formats. For context, netCDF is a binary format that offers the advantage to be compact and efficient when dealing with multidimensional data, while the others have the advantage of being human-readable (but can be very inefficient for large datasets). None of these formats is optimized for cloud object storage, although there are ongoing efforts for netCDF (e.g. Zarr and Icechunk). For this reason, Parquet has been drawing more and more attention from the earth sciences community recently.

Parquet is a cloud-optimized binary format for tidy and large data (i.e. large tables) that is

language-agnostic. It is widely used in the data science and corporate worlds, and the software ecosystem around it is sound, mature, and growing. Table 1 presents an overview of the characteristics of Parquet and the other formats. We chose Parquet as the target format for CrocoLake because (1) it is optimized for cloud storage and cloud computing, (2) its mature software ecosystem includes packages in multiple coding languages to access it (Python, Julia, MATLAB, web technologies, etc.), and (3) novel users are generally more familiar with tabular data than with specialized oceanographic data formats.

As we aim to make CrocoLake easily accessible from a technical standpoint, the main drawbacks of Parquet at present are that (1) many workflows in ocean modeling are based on multidimensional data structures (not tabular ones) and (2) attaching attributes to the data is not as straightforward as it could be. However, we see CrocoLake as a major building bloc in workflows that require fast access to sparse in-situ ocean observations, responding to necessities of data storage with fast access both on disk and the cloud, rather than as an end-all solution for ocean observations. For example, array-based storage formats might instead be more relevant to workflows that rely mostly on regular and gridded observations (such as satellite recordings).

Table 1. Comparison between different common file formats for oceanographic datasets.

Feature Name	CSV	netCDF	Parquet	ASCII	JSON	Zarr + Icechunk
Cloud-Optimized Structure	No Tabular	No Array-based	Yes Tabular	No Tabular	No Hierarchical	Yes Array-based
Available tools in:						
Python	Yes	Yes	Yes	Yes	Yes	Yes
Julia	Yes	Yes	Yes	Yes	Yes	Zarr only
MATLAB	Yes	Yes	Yes	Yes	Yes	Zarr only
Attributes descriptors	Dedicated columns, header	Dictionary, accessed with data	Dictionary, separate access from data	Dedicated column	Dedicated field in object	Dictionary, accessed with data

Another key feature of CrocoLakeTools is that, unlike most aforementioned projects, it is fully open-source. Anyone can thus use CrocoLakeTools to build their own flavor of CrocoLake. We also welcome new contributors to CrocoLakeTools, for example by adding converters to support new datasets.

Code architecture

Converters

The core task of CrocoLakeTools is to take one or more files from a dataset and convert them to parquet, ensuring that CrocoLake's schema is followed. This is achieved through the methods contained in the Converter class and its subclasses. While the conversion of all datasets requires some general functionality (e.g. renaming the original variables to the final schema), each conversion requires specialized tools that differ between datasets (e.g. the map used to rename variables). CrocoLakeTools then contains the Converter class, which contains the methods shared across datasets, and from which converter subclasses inherit and implement the specific needs of each dataset.

Workflow

Local mirrors

The first step in our workflow is to retrieve original files from each data provider (Figure 1). The original source data follow the format, schema, nomenclature, and conventions defined by their side (project, mission, scientist, etc.) independently of CrocoLake's workflow. Modules to download the original data are optional. They are expected to inherit from the Downloader class and be called `downloader<DatasetName>` (e.g. `downloaderArgoGDAC`). At the time of writing CrocoLakeTools is released with a downloader to build a local mirror of the Argo Global Data Assembly Center (GDAC), and we hope to support more data providers in the future. Whether a downloader module exists or the user downloads the data themselves, the original data is stored on disk and this is the starting point for the converter.

Parquet datasets

The second step is to convert the data to parquet, and finally merge the datasets into CrocoLake (Figure 2). The core of CrocoLakeTools are the modules in the Converter class and its subclasses. Each original dataset has its own subclass called `converter<DatasetName>`, e.g. `converterGLODAP`; further specifiers can be added as necessary (e.g. at this time there are a few different converters for Argo data to prepare different datasets). The need for a dedicated converter for each project despite the usage of common data formats (e.g. netCDF, CSV) is due to differences in schema (e.g. variable names or units). Depending on the dataset, multiple converters can be applied. For example, to create CrocoLake, Argo data goes through two converters: 1. `converterArgoGDAC`, which converts the original Argo GDAC preserving most of its original conventions; 2. `converterArgoQC`, which takes the output of the previous step and applies some filtering based on Argo's QC flags and makes the data conforming to CrocoLake's schema. At this time CrocoLakeTools is released with converters for data from Argo (Wong et al., 2020), Spray Data (Sherman et al., 2002), (Rudnick et al., 2016) and GLODAP (Global Ocean Data Analysis Project) (Lauvset et al., 2016), (Olsen et al., 2016).

CrocoLake

CrocoLake is one parquet dataset that contains all converted datasets merged together. This can be achieved with the script `merge_crocolake.py`. The script first creates a directory containing symbolic links to each converted dataset. It then uses the submodule CrocoLakeLoader to seamlessly load all the converted datasets into memory as one dask dataframe with a uniform schema, merges them into CrocoLake, and stores it back to disk.

CrocoLake can be accessed with several programming languages with just a few lines of codes: the submodules CrocoLake-Python, CrocoLake-Matlab, and CrocoLake-Julia contain tools and examples in some languages.




Figure 1: CrocoLake's workflow: 'downloader's. 'CrocoLakeTools' is set up to host modules that are dedicated to download the desired datasets from the web. It currently supports the download only of Argo data (solid line subset), and other datasets require the user to download them manually (dashed border subsets). Argo, Spray Data and GLODAP (Global Ocean Data Analysis Project) are the different data providers supported at the time of submission.




Figure 2: CrocoLake's workflow. 'converter's read the data in their original format, transform it following CrocoLake's conventions, converts it to parquet, and stores it back to disk. Each dataset is converted to its own parquet version. Thanks to the submodule 'CrocoLakeLoader', multiple parquet datasets are merged into a uniform dataframe which is saved to disk as CrocoLake. For each dataset, a version containing only physical variables ('PHY') or also biogeochemical variables ('BGC') can be generated.

155 Schema

156 The nomenclature, units and data types are generally based on Argo's (<https://vocab.nerc.ac.uk/collection/R03/current/>). For CrocoLake's variables that are not present in the
157 Argo program, we provide new names mantaining consistency with Argo's style.
158

159 Profile numbering

160 Ocean data is often accessed by profiles and we provide this functionality for CrocoLake too:
161 the user can retrieve the profiles through the CYCLE_NUMBER variable, which is unique and
162 progressive for each PLATFORM_NUMBER of each subdataset (DB_NAME). As each original product
163 uses its own conventions, the CYCLE_NUMBER of some datasets is generated ad hoc during their
164 conversion if no obvious match with CYCLE_NUMBER exists. The procedure for each dataset is
165 detailed in the online documentation.

166 Quality control

167 CrocoLake contains only quality-controlled (QC) measurements. We rely exclusively on QCs
168 performed by the data providers and at the time we do not perform any additional QC ourselves
169 (although this might change in the future). Each parameter <PARAM> has a corresponding
170 <PARAM>_QC flag that is generally set to 1 to indicate that the data is reliable. For Argo
171 measurements, the original QC value is preserved, and only measurements with QC values of
172 1, 2, 5, and 8 considered.

173 Measurement errors

174 Each parameter <PARAM> has a corresponding <PARAM>_ERROR that indicates a measurement's
175 error as provided in the original dataset. When no error is provided, <PARAM>_ERROR is set to
176 null.

177 Documentation and updates

178 The [documentation](#) describes the specifics of each dataset (e.g. what quality-control filter we
179 apply to each dataset, the procedure to generate the profile numbers, etc.), and get updated
180 every time a new feature is made available.

181 Citation

182 If you use CrocoLakeTools and/or CrocoLake, please do not limit yourself to citing this
183 manuscript but also remember to cite the datasets that you have used as indicated in the
184 documentation. For example, if your work relies on Argo measurements, acknowledge Argo
185 ([Wong et al., 2020](#)). This is important both for the maintainers of each product to track their
186 impact and to acknowledge their efforts that made your work possible.

187 Acknowledgements

188 [NASA ECCO NSC...]. This material is based upon work supported by the National Science
189 Foundation under Grant No. 2311382, "Collaborative Research: Frameworks: A community
190 platform for accelerating observationally-constrained regional oceanographic modeling." The
191 Argo data are collected and made freely available by the International Argo Program and the
192 national programs that contribute to it (<https://argo.ucsd.edu>, <https://www.ocean-ops.org>).
193 The Argo Program is part of the Global Ocean Observing System.

References

- Forget, Gael. (2025). *Euroargodev/ArgoData.jl* [Data set]. Zenodo. <https://doi.org/https://doi.org/10.5281/zenodo.3633717>
- Forget, Gaël, & Gebbie, G. J. (2025). *JuliaOcean/OceanRobots.jl* [Data set]. Zenodo. <https://doi.org/https://doi.org/10.5281/zenodo.4718472>
- Kelley, D. E., Harbin, J., & Richards, C. (2021). argoFloats: An r package for analyzing argo data. *Frontiers in Marine Science*, 8, 635922.
- Kelley, D. E., Richards, C., & Layton, C. (2022). Oce: An r package for oceanographic analysis. *Journal of Open Source Software*, 7(71), 3594.
- Lauvset, S. K., Key, R. M., Olsen, A., Van Heuven, S., Velo, A., Lin, X., Schirnack, C., Kozyr, A., Tanhua, T., Hoppema, M., & others. (2016). A new global interior ocean mapped climatology: The 1× 1 GLODAP version 2. *Earth System Science Data*, 8(2), 325–340.
- Maze, G., & Balem, K. (2020). Argopy: A python library for argo ocean data analysis. *Journal of Open Source Software*, 5.
- Mishonov, A. V., Boyer, T. P., Baranova, O. K., Bouchard, C. N., Cross, S. L., Garcia, H. E., Locarnini, R. A., Paver, C. R., Reagan, J. R., Wang, Z., & others. (2024). *World ocean database 2023*.
- Olsen, A., Key, R. M., Van Heuven, S., Lauvset, S. K., Velo, A., Lin, X., Schirnack, C., Kozyr, A., Tanhua, T., Hoppema, M., & others. (2016). The global ocean data analysis project version 2 (GLODAPv2)—an internally consistent data product for the world ocean. *Earth System Science Data*, 8(2), 297–323.
- Rudnick, D. L., Davis, R. E., & Sherman, J. T. (2016). Spray underwater glider operations. *Journal of Atmospheric and Oceanic Technology*, 33(6), 1113–1122.
- Sherman, J., Davis, R. E., Owens, W., & Valdes, J. (2002). The autonomous underwater glider" spray". *IEEE Journal of Oceanic Engineering*, 26(4), 437–446.
- Szekely, T., Gourrion, J., Pouliquen, S., Reverdin, G., & Merceur, F. (2019). *CORA, coriolis ocean dataset for reanalysis*.
- Team, T. I. (2018). *International quality controlled ocean database (IQuOD) version 0.1 - aggregated and community quality controlled ocean profile data 1772-2018 (NCEI accession 0170893)* [Data set]. NOAA National Centers for Environmental Information. <https://doi.org/https://doi.org/10.7289/v51r6nsf>
- Tucker, T., Giglio, D., Scanderbeg, M., & Shen, S. S. (2020). Argovis: A web application for fast delivery, visualization, and analysis of argo data. *Journal of Atmospheric and Oceanic Technology*, 37(3), 401–416.
- Wong, A. P., Wijffels, S. E., Riser, S. C., Pouliquen, S., Hosoda, S., Roemmich, D., Gilson, J., Johnson, G. C., Martini, K., Murphy, D. J., & others. (2020). Argo data 1999–2019: Two million temperature-salinity profiles and subsurface velocity observations from a global array of profiling floats. *Frontiers in Marine Science*, 7, 700.