The Journal of Open Source Software

CrocoLakeTools: A Python package to convert ocean observations to the parquet format

Enrico Milanese [1], **David Nicholson** [1], **Gaël Forget**[2], **and Susan Wijffels**[1]

1 Woods Hole Oceanographic Institution, United States of America **2** Massachusetts Institute of Technology, United States of America

DOI: 10.xxxxxx/draft

Software

- Review
- Repository
- Archive

Editor: Open Journals

Reviewers:

- @openjournals

Submitted: 01 January 1970
Published: unpublished

Summary

Investigations of the ocean state are possible thanks to the ever growing number of measurements performed with multiple instruments by different research missions. The vast and heterogeneous efforts have brought the community to define data storage conventions (e.g. CF-netCDF) and to assemble collections of datasets (e.g. the World Ocean Database). Yet, accessing these datasets often requires the usage of specialized tools, is generally inefficient over the cloud, and remains a major obstacle to new users in terms of pre-required knowledge and time. CrocoLakeTools is a Python package that addresses those challenges by providing workflows to convert various oceanographic datasets from their original format to a uniform parquet dataset with a shared schema.

Statement of need

`CrocoLakeTools` is a Python package to build workflows that convert ocean observations from their original formats (e.g. netCDF, CSV) to parquet. CrocoLakeTools takes advantage of Python's well-established and growing ecosystem of open tools: it uses `dask`'s parallel computing capabilities to convert multiple files at once and to handle larger-than-memory data. `dask` is already well-integrated with `xarray` and `pandas`, two widely used Python libraries for the treatment of array and tabular data, respectively, and with `pyarrow`, the API to the Apache Arrow library which is used to generate the parquet dataset.

Parquet is a data storage format for big tidy data which presents several advantages: it is language agnostic (it can be accessed with Python, Matlab, Julia and web developement technologies); it offers faster reading performances than other tabular formats such as CSV; it is optimized for cloud systems storage and operations; it is widespread in the data science community, leading to a multitude of freely accessible tools and educational material.

`CrocoLakeTools` was developed with the goal of building and serving CrocoLake, a regularly refreshed database of sparse in-situ oceanographic observations that are pre-filtered to contain only quality-controlled measurements. `CrocoLakeTools` was designed to be used by researchers, engineers and data scientists in oceanography and to be accessed by the wider oceanographic community.

State of the field

In-situ measurements of ocean characteristics are collected by a wide range of projects at different scales: from small research groups to international consortia, from time-limited regional explorations to continuous global observation arrays. This has led to a highly heterogeneous landscape of products that often differ for schema (variable names, data types), data format,

and access services. As a result, ingesting and analyzing data from different sources, e.g. for research purposes, presents technical and logistical challenges. The need for homogenized ocean data products has then prompted several efforts, in particular when it comes to sparse observations, which are characterized by discrete measurements that are generally irregular and sparse in time and space.

One of the most comprehensive projects to date is the World Ocean Database (WOD) (Mishonov et al., 2024), which contains more than 40 variables gathered from several sources (buoys, gliders, floats, bathythermographs, etc.). The earliest record dates back to 1772, the data is quality-controlled, and the database is regularly updated with the latest measurements available. The data is public and freely available through different interfaces (WODselect, THREDDS, HTTPS, FTP) in ASCII, Comma Separated Value (CSV) and netCDF formats. WOD is maintained by the National Centers for Environmental Information (NCEI) of the United States' National Oceanic and Atmospheric Administration (NOAA).

Copernicus Marine Service is an initiative of the European Union that provides a wide range of ocean data products: physical, biogeochemical and sea ice data at regional and global scales, obtained from numerical models, in-situ observations and satellite observations. Among Copernicus' products, CORA(Szekely et al., 2019) (Coriolis Ocean database for ReAnalysis) provides quality-controlled in-situ temperature and salinity measurements gathered from several global and regional networks.

The International Quality-controlled Ocean Database (IQuOD) (Team, 2018) is an effort by the oceanographic community "with the goal of producing the highest quality and complete single ocean profile repository along with (intelligent) metadata and assigned uncertainties". It includes subsurface ocean profiles of several variables, paying particular attention to temperature measurements. The database is prepared by NCEI, and it is freely available in netCDF format through multiple channels (THREDDS, HTTPS, FTP).

The Ocean Data Platform by the non-profit HUB Ocean is among the youngest projects in the community, and it allows to find and access datasets from a catalog. The user can interact with the platform through different interfaces (SDK, REST API, OQS, JupyterHub workspaces), loading the datasets in tabular format.

Of particular interests to the community has been the data collected by the Argo (Wong et al., 2020) international program with a growing fleet of robotic floats, and several tools have been developed to handle such data. argopy (Maze & Balem, 2020) is a Python library that facilitates access and manipulation of data from the real-time observing system Argo (Wong et al., 2020). The user provides some filters (e.g. coordinate and time ranges, measurement name, etc) and argopy retrieves the relevant data from the Argo Global Data Assembly Center (GDAC). ArgoData.jl(Gael Forget, 2025) and OceanRobots.jl(Gaël Forget & Gebbie, 2025) are Julia packages that provide similar functionalities to argopy. Argovis (Tucker et al., 2020) is a REST API and web application hosted at the University of Colorado, Boulder (United States) and serving profile data in JSON format from the Argo program, and from ship and drifters missions. oce (Kelley et al., 2022) is a package written in R providing multiple functions to access oceanographic observations stored in a variety of formats. It focuses on datasets for physical oceanography with the goal of providing more tools in the future. argoFloats (Kelley et al., 2021) is a derived package that focuses on fetching and manipulating Argo data.

The above efforts often serve data in ASCII, CSV, JSON, or netCDF formats. For context, netCDF is a binary format that offers the advantage to be compact and efficient when dealing with multidimensional data, while the others have the advantage of being human-readable (but can be very inefficient for large datasets). None of these formats is optimized for cloud object storage, although there are ongoing efforts for netCDF (e.g. Zarr and Icechunk). For this reason, Parquet has been drawing more and more attention from the earth sciences community recently.

Parquet is a cloud-optimized binary format for tidy and large data (i.e. large tables) that is

language-agnostic. It is widely used in the data science and corporate worlds, and the software ecosystem around it is sound, mature, and growing. Table 1 presents an overview of the characteristics of Parquet and the other formats. We chose Parquet as the target format for `CrocoLake` because (1) it is optimized for cloud storage and cloud computing, (2) its mature software ecosystem includes packages in multiple coding languages to access it (Python, Julia, MATLAB, web technologies, etc.), (3) its tabular format suits in-situ observations which are sparse and do not benefit from advantages of array-based formats as much as gridded datasets, and (4) novel users are generally more familiar with tabular data than with specialized oceanographic data formats.

As we aim to make CrocoLake easily accessible from a technical standpoint, the main drawbacks of Parquet at present are that (1) many workflows in ocean modeling are based on multidimensional data structures (not tabular ones) and (2) attaching attributes to the data is not straightfoward (compared to netCDF). However, we see CrocoLake as a major building bloc in workflows that require fast access to sparse in-situ ocean observations, responding to necessities of data storage with fast access both on disk and the cloud, rather than as an end-all solution for ocean observations. For example, array-based storage formats might instead be more relevant to workflows that rely mostly on regular and gridded observations (such as satellite recordings).

Table 1. Comparison between different common file formats for oceanographic datasets.

Feature Name	CSV	netCDF	Parquet	ASCII	JSON	Zarr + Icechunk
Cloud-Optimized	No	No	Yes	No	No	Yes
Structure	Tabular	Array-based	Tabular	Tabular	Hierarchical	Array-based
Available tools in:						
Python	Yes	Yes	Yes	Yes	Yes	Yes
Julia	Yes	Yes	Yes	Yes	Yes	Zarr only
MATLAB	Yes	Yes	Yes	Yes	Yes	Zarr only
Attributes descriptors	Dedicated columns, header	Dictionary, accessed with data	Dictionary, separate access from data	Dedicated column	Dedicated field in object	Dictionary, accessed with data

Another key feature of CrocoLakeTools is that it is fully open-source. Anyone can thus use CrocoLakeTools to build their own flavor of CrocoLake. We also welcome new contributors to CrocoLakeTools, for example by adding converters to support new datasets.

Code architecture

Converters

The core task of CrocoLakeTools is to take one or more files from a dataset and convert them to parquet, ensuring that CrocoLake's schema is followed. This is achieved through the methods contained in the `Converter` class and its subclasses. While the conversion of all datasets requires some general functionality (e.g. renaming the original variables to the final schema), each conversion requires specialized tools that differ between datasets (e.g. the map used to rename variables). CrocoLakeTools then contains the `Converter` class, which contains the methods shared across datasets, and from which converter subclasses inherit and implement the specific needs of each dataset.

Workflow

Local mirrors

The first step in our workflow is to retrieve original files from each data provider (Figure 1). The original source data follow the format, schema, nomenclature, and conventions defined by their provider (project, mission, scientist, etc.) independently of CrocoLake's workflow. Modules to download the original data are optional. They are expected to inherit from the `Downloader` class and be called `downloader<DatasetName>` (e.g. `downloaderArgoGDAC`). At the time of writing CrocoLakeTools is released with a downloader to build a local mirror of the Argo Global Data Assembly Center (GDAC), and we hope to support more data providers in the future. Whether a downloader module exists or the user downloads the data themselves, the original data is stored on disk and this is the starting point for the converter.

Parquet datasets

The second step is to convert the data to parquet, and finally merge the datasets into CrocoLake (Figure 2). The core of `CrocoLakeTools` are the modules in the `Converter` class and its subclasses. Each original dataset has its own subclass called `converter<DatasetName>`, e.g. `converterGLODAP`; further specifiers can be added as necessary (e.g. at this time there a few different converters for Argo data to prepare different datasets). The need for a dedicated converter for each project despite the usage of common data formats (e.g. netCDF, CSV) is due to differences in schema (e.g. variable names or units). Depending on the dataset, multiple converters can be applied. For example, to create CrocoLake, Argo data goes through two converters: 1. `converterArgoGDAC`, which converts the original Argo GDAC preserving most of its original conventions; 2. `converterArgoQC`, which takes the output of the previous step and applies some filtering based on Argo's QC flags and makes the data conforming to CrocoLake's schema. At this time CrocoLakeTools is released with converters for data from Argo (Wong et al., 2020), GLODAP (Global Ocean Data Analysis Project, Key et al. (2015), Lauvset et al. (2016), Olsen et al. (2016)) and Spray Data (Sherman et al. (2002), Rudnick et al. (2004), Rudnick & Cole (2011), Gopalakrishnan et al. (2013), Johnston et al. (2013), Rudnick et al. (2013), Rudnick et al. (2015), Schönau & Rudnick (2015), Rudnick (2016), Rudnick et al. (2016), Todd et al. (2016), Centurioni et al. (2017), Schönau & Rudnick (2017), Zeiden et al. (2019), Heiderich & Todd (2020)). The choice of glider data is limited to Spray Data because they are provided in delayed mode (i.e. with comprehensive and manual quality-control): we are interested in including data from the wider IOOS gliders network, yet to our knowledge quality-control processes for delayed mode data are not yet standardized across all the network. We are also working towards adding new datasets, such as observations collected by Saildrones used by the Ocean Climate Stations Project and expendable bathythermograph data collected by the Oleander Ship Of Opportunity project.

CrocoLake

CrocoLake is one parquet dataset that contains all converted datasets merged together. This can be achieved with the script `merge_crocolake.py`. The script first creates a directory containing symbolic links to each converted dataset. It then uses the submodule `CrocoLakeLoader` to seamlessly load all the converted datasets into memory as one dask dataframe with a uniform schema, merges them into CrocoLake, and stores it back to disk.

CrocoLake can be accessed with several programming languages with just a few lines of codes: the submodules CrocoLake-Python, CrocoLake-Matlab, and CrocoLake-Julia contain tools and examples in some languages.

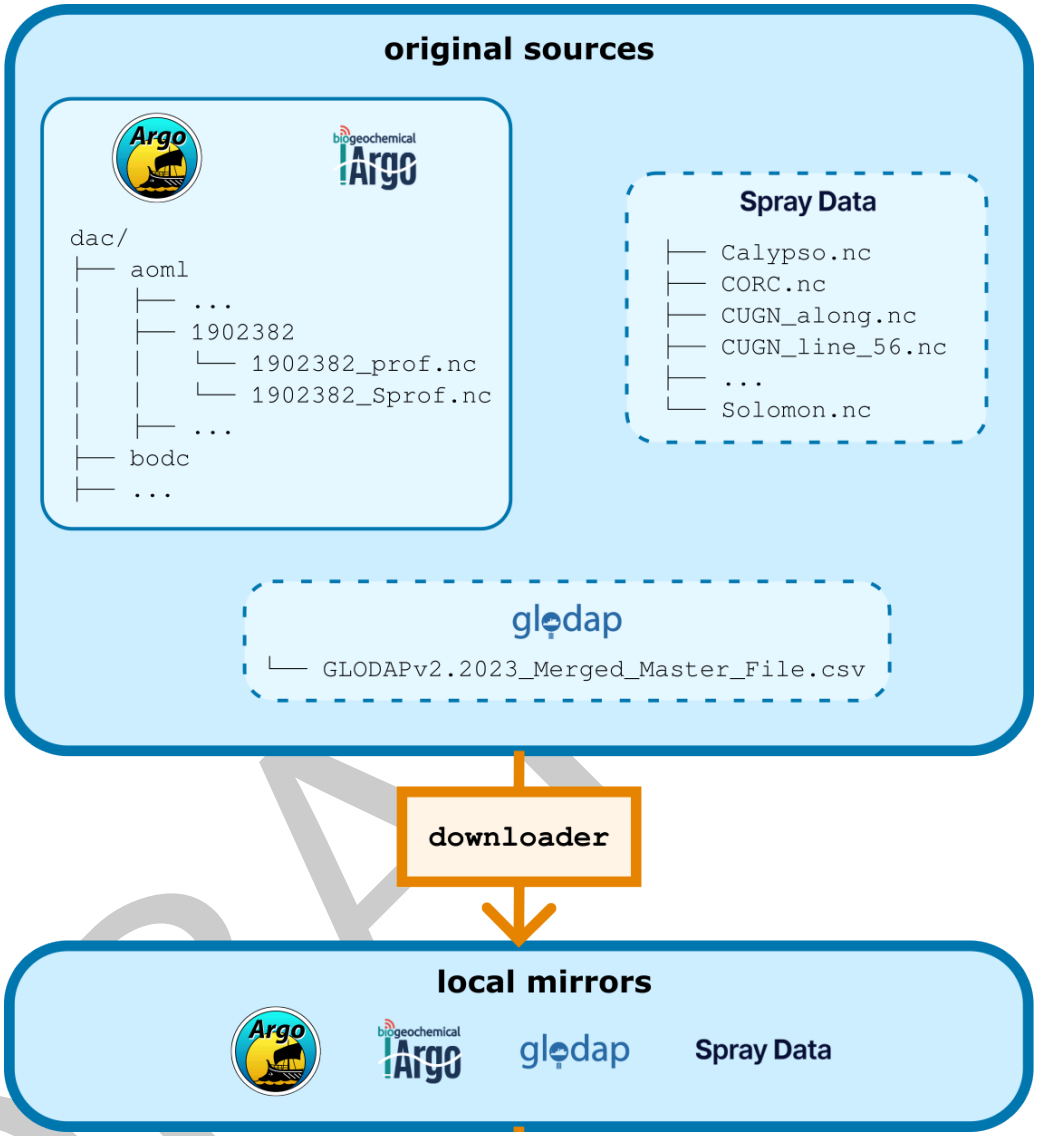

Figure 1: CrocoLake's workflow: 'downloader's. 'CrocoLakeTools' is set up to host modules that are dedicated to download the desired datasets from the web. It currently supports the download only of Argo data (solid line subset), and other datasets require the user to download them manually (dashed border subsets). Argo, Spray Data and GLODAP (Global Ocean Data Analysis Project) are the different data providers supported at the time of submission.

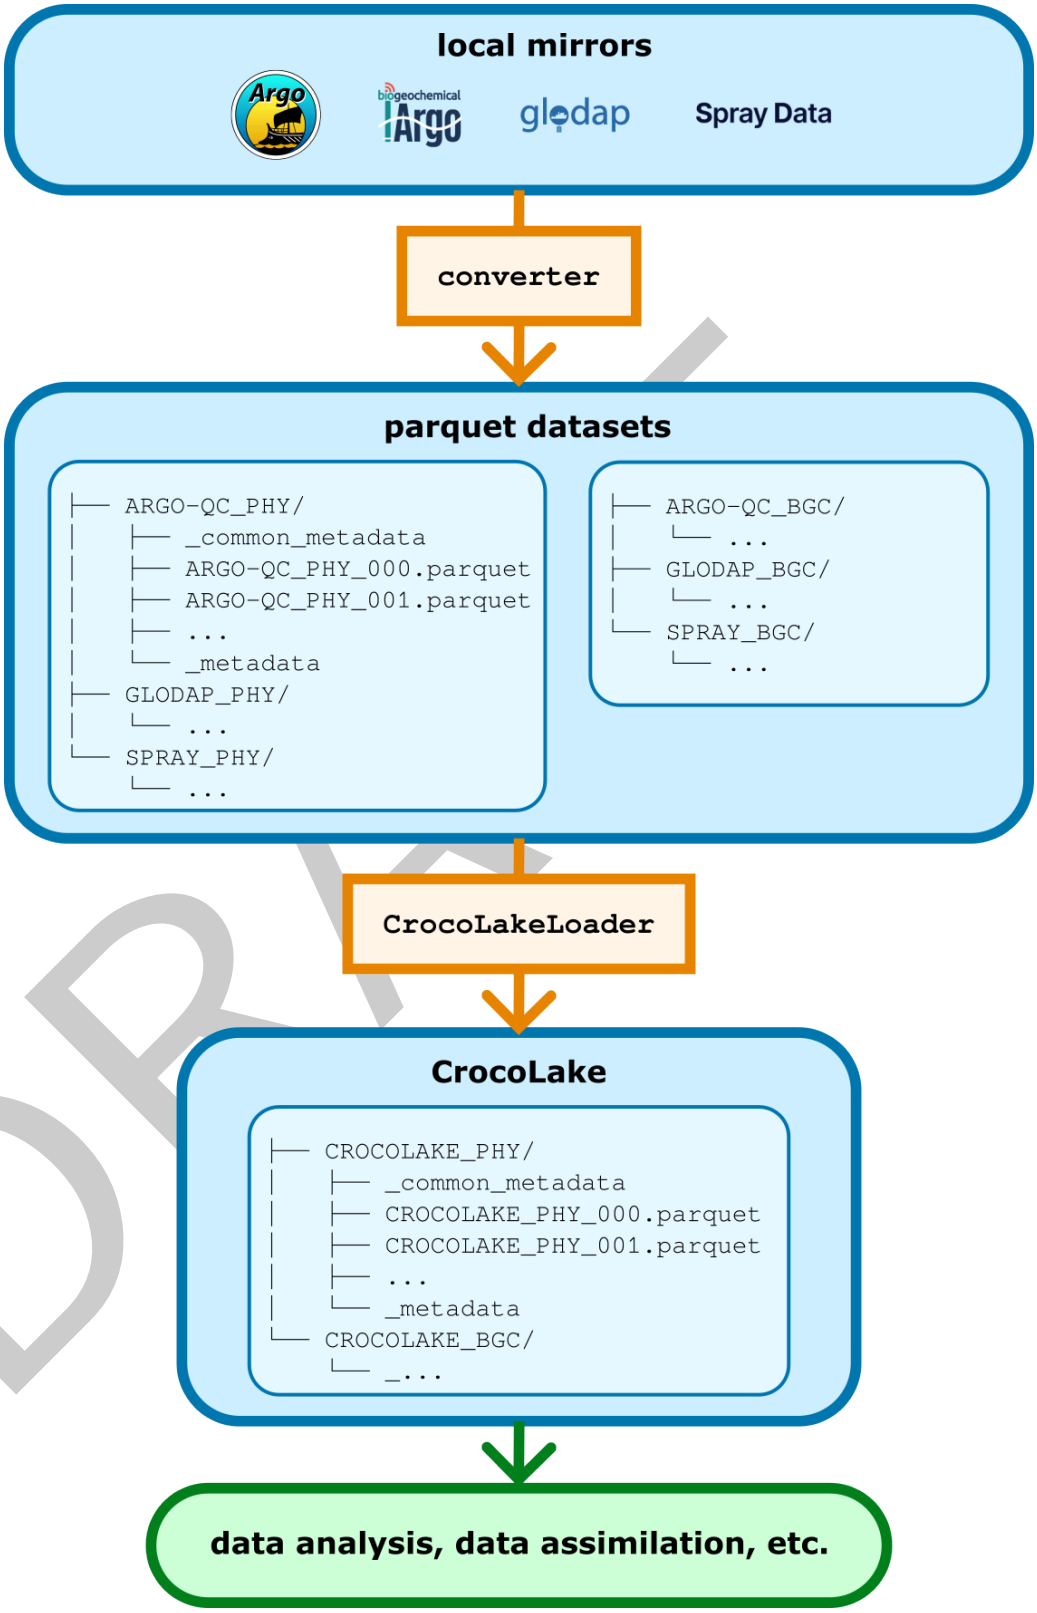

Figure 2: CrocoLake's workflow: 'converter's. 'converter's read the data in their original format, transform it following CrocoLake's conventions, converts it to parquet, and stores it back to disk. Each dataset is converted to its own parquet version. Thanks to the submodule 'CrocoLakeLoader', multiple parquet datasets are merged into a uniform dataframe which is saved to disk as CrocoLake. For each dataset, a version containing only physical variables ('PHY') or also biogeochemical variables ('BGC') can be generated.

Schema and metadata

The nomenclature, units and data types are generally based on Argo's (https://vocab.nerc.ac.uk/collection/ and are detailed in the online documentation. For CrocoLake's variables that are not present in the Argo program, we provide new names maintaining consistency with Argo's style.

The parquet format has two different types of metadata: metadata related to storage (e.g. schema, number of rows, etc.) and column attributes. We use the latter type to store information about the units for each variable present in CrocoLake. At the moment we are not including other metadata from the original sources (e.g. information about principal investigator, contact details, comments, etc.), but we are exploring solutions to include some of this in the future, based also on feedback from users.

Profile numbering

Ocean data is often accessed by profiles and we provide this functionality for CrocoLake too: the user can retrieve the profiles through the `CYCLE_NUMBER` variable, which is unique and progressive for each `PLATFORM_NUMBER` of each subdataset (`DB_NAME`). As each original product uses its own conventions, the `CYCLE_NUMBER` of some datasets is generated ad hoc during their conversion if no obvious match with `CYCLE_NUMBER` exists. The procedure for each dataset is detailed in the online documentation.

Quality control

CrocoLake contains only quality-controlled (QC) measurements. We rely exclusively on QCs performed by the data providers and at the time we do not perform any additional QC ourselves (although this might change in the future). Each parameter `<PARAM>` has a corresponding `<PARAM>_QC` flag that is generally set to 1 to indicate that the data is reliable. For Argo measurements, the original QC value is preserved, and only measurements with QC values of 1, 2, 5, and 8 considered.

Measurement errors

Each parameter `<PARAM>` has a corresponding `<PARAM>_ERROR` that indicates a measurement's error as provided in the original dataset. When no error is provided, `<PARAM>_ERROR` is set to null.

Benchmarking

We here present a few examples of the advantages and disadvantages in accessing the generated parquet datasets. We compare access to Argo data in a few different scenarios between CrocoLake and `argopy` (Maze & Balem, 2020), which is widely used in the community. For the comparison, we use the parquet Argo data generated by `converterArgoGDAC` (i.e. without applying any filtering based on quality-control flags), as it is the version closest to the data served by `argopy`. The parquet dataset is stored on the cloud in a Amazon Web Services Simple Storage Solution (AWS S3) bucket, and `argopy` is called to access data from the Argo GDAC, so the data is always accessed over the internet. The parquet dataset is accessed with Python for comparison with `argopy`; we use `dask` but other packages exist to access parquet data (e.g. `duckDB`, `polars`) and the choice of package can affect the performance.

The goal of this benchmarking exercise is to provide a ballpark estimate of reading performances in a few different scenarios to showcase the potential of including oceanographic datasets in parquet format in scientific workflows; comparison with an existing tool as `argopy` is qualitative to provide a reference to help evaluate the performance of the parquet-based workflow. More generally, the reading performance of a given dataset depends on multiple factors beyond the dataset format itself (coding language, selection parameters, local vs cloud-based access, network connectivty, storage service, etc.).

Benchmarks were performed on a Lenovo ThinkPad P14s Gen 4 equipped with an AMD Ryzen 7 PRO 7840U processor (8 cores, 16 threads, up to 5.13 GHz) and 26 GiB DDR4 RAM. Network connectivity was provided via Wi-Fi: bandwidth measurements indicated an average download speed of 466.41 Mbit/s (10 runs, range: 363.37–555.70 Mbit/s) and an upload speed of 27.34 Mbit/s (10 runs, range: 26.98–27.48 Mbit/s); latency to google.com averaged 26.2 ms (10 measurements, range: 20.9–27.6 ms). All non-essential applications were closed during benchmarking, and measurements were collected immediately prior to the tests to ensure reproducibility.

The following Tables 2 and 3 summarize data loading times for PHY and BGC Argo datasets from AWS S3 (Parquet) versus the `argopy` Python library. Benchmarks include space and time filtering and loading by float WMO ID.

Table 2 – North West Atlantic Regional Subsets (2017)

Data type	Subset size	Filter details	Parquet time	`argopy` time
PHY	Large	Lat: [25,60], Lon: [-80,-30], Date: 2017-01-01 to 2018-01-01	44.8 ± 0.4 s	1266 ± 8.6s
PHY	Small	Lat: [55,60], Lon: [-50,-55], Date: 2017-01-01 to 2018-01-01	19.5 ± 1.5 s	63 ± 1.9s
BGC	Large	Lat: [25,60], Lon: [-80,-30], Date: 2017-01-01 to 2018-01-01	8.5 ± 0.7 s	954 ± 1.1s
BGC	Small	Lat: [55,60], Lon: [-50,-55], Date: 2017-01-01 to 2018-01-01	5.0 ± 2.3 s	108 ± 4.3s

Table 3 – Loading Full Float Records

Data type	Floats	Parquet time	`argopy` time
PHY	1 float	4.2 ± 1.3 s	3.0 ± 0.3 s
PHY	20 floats	17.4 ± 0.8 s	22.5 ± 1.1 s
BGC	1 float	3.4 ± 1.8 s	6.2 ± 0.6 s
BGC	20 floats	8.0 ± 2.8 s	140.0 ± 6.0 s

Notes:

- All timings are mean ± standard deviation measured over ten runs, except for the "Large" subsets in Table 2 which are measured over two runs.
- Large/small regional subsets (Table 2) differ by latitude/longitude bounding boxes (see above).
- The floats of the records in Table 3 were randomly picked and are identified by the following WMO IDs:
  - PHY, single float: 6902801;
  - PHY, 20 floats: 4902120, 6902707, 3901629, 6901177, 4901812, 4901765, 4901216, 4902322, 6902801, 6902753, 6902805, 5904749, 4901827, 4902419, 4902397, 4901755, 6902810, 5904007, 4902112, 1901596;
  - BGC, single float: 4902410;

- BGC, 20 floats: 6901754, 6901030, 4902409, 6902810, 6901750, 6902805, 6901182, 6902686, 5904989, 6901603, 4902390, 5904770, 1901217, 6901751, 1901208, 6901752, 6902812, 6901023, 6901758, 6901524.

Documentation and updates

The documentation describes the specifics of each dataset (e.g. what quality-control filter we apply to each dataset, the procedure to generate the profile numbers, etc.), and is updated every time a new feature is made available.

This document reflects the code at the time of publication. Any future updates and changes to the code will be documented in the code `main` branch, which will contain the most recent and best version available of `CrocoLakeTools`.

Citation

If you use CrocoLakeTools and/or CrocoLake, please do not limit yourself to citing this manuscript but also remember to cite the datasets that you have used as indicated in the documentation. For example, if your work relies on Argo measurements, acknowledge Argo (Wong et al., 2020). This is important both for the maintainers of each product to track their impact and to acknowledge their efforts that made your work possible.

Acknowledgements

This material is based upon work supported by the National Science Foundation under Grant No. 2311382, "Collaborative Research: Frameworks: A community platform for accelerating observationally-constrained regional oceanographic odeling." G.F. acknowledges funding from NASA award 1686358, NOAA award 035724-00001, and ARIA's POLEMIX award.

We acknowledge the Argo, Spray Data, and GLODAP programs for their efforts in collecting ocean observations and making them publicly available. The Argo data are collected and made freely available by the International Argo Program and the national programs that contribute to it (https://argo.ucsd.edu, https://www.ocean-ops.org), and is part of the Global Ocean Observing System. E.M. thanks Robert Todd for guidance in accessing Spray Data.

References

Centurioni, L. R., Hormann, V., Talley, L. D., Arzeno, I., Beal, L., Caruso, M., Conry, P., Echols, R., Fernando, H. J., Giddings, S. N., & others. (2017). Northern arabian sea circulation-autonomous research (NASCar) a research initiative based on autonomous sensors. *Oceanography*, *30*(2), 74–87.

Forget, Gael. (2025). *Euroargodev/ArgoData.jl* [Data set]. Zenodo. https://doi.org/https://doi.org/10.5281/zenodo.3633717

Forget, Gaël, & Gebbie, G. J. (2025). *JuliaOcean/OceanRobots.jl* [Data set]. Zenodo. https://doi.org/https://doi.org/10.5281/zenodo.4718472

Gopalakrishnan, G., Cornuelle, B. D., Hoteit, I., Rudnick, D. L., & Owens, W. B. (2013). State estimates and forecasts of the loop current in the gulf of mexico using the MITgcm and its adjoint. *Journal of Geophysical Research: Oceans*, *118*(7), 3292–3314.

Heiderich, J., & Todd, R. E. (2020). Along-stream evolution of gulf stream volume transport. *Journal of Physical Oceanography*, *50*(8), 2251–2270.

Johnston, T. S., Rudnick, D. L., Alford, M. H., Pickering, A., & Simmons, H. L. (2013). Internal tidal energy fluxes in the south china sea from density and velocity measurements

by gliders. *Journal of Geophysical Research: Oceans*, *118*(8), 3939–3949.

Kelley, D. E., Harbin, J., & Richards, C. (2021). argoFloats: An r package for analyzing argo data. *Frontiers in Marine Science*, *8*, 635922.

Kelley, D. E., Richards, C., & Layton, C. (2022). Oce: An r package for oceanographic analysis. *Journal of Open Source Software*, *7*(71), 3594.

Key, R. M., Olsen, A., Heuven, S. van, Lauvset, S. K., Velo, A., Lin, X., Schirnick, C., Kozyr, A., Tanhua, T., Hoppema, M., & others. (2015). Global ocean data analysis project, version 2 (GLODAPv2). *Ornl/Cdiac-162, Ndp-093*.

Lauvset, S. K., Key, R. M., Olsen, A., Van Heuven, S., Velo, A., Lin, X., Schirnick, C., Kozyr, A., Tanhua, T., Hoppema, M., & others. (2016). A new global interior ocean mapped climatology: The 1× 1 GLODAP version 2. *Earth System Science Data*, *8*(2), 325–340.

Maze, G., & Balem, K. (2020). Argopy: A python library for argo ocean data analysis. *Journal of Open Source Software*, *5*.

Mishonov, A. V., Boyer, T. P., Baranova, O. K., Bouchard, C. N., Cross, S. L., Garcia, H. E., Locarnini, R. A., Paver, C. R., Reagan, J. R., Wang, Z., & others. (2024). *World ocean database 2023*.

Olsen, A., Key, R. M., Van Heuven, S., Lauvset, S. K., Velo, A., Lin, X., Schirnick, C., Kozyr, A., Tanhua, T., Hoppema, M., & others. (2016). The global ocean data analysis project version 2 (GLODAPv2)–an internally consistent data product for the world ocean. *Earth System Science Data*, *8*(2), 297–323.

Rudnick, D. L. (2016). Ocean research enabled by underwater gliders. *Annual Review of Marine Science*, *8*, 519–541.

Rudnick, D. L., & Cole, S. T. (2011). On sampling the ocean using underwater gliders. *Journal of Geophysical Research: Oceans*, *116*(C8).

Rudnick, D. L., Davis, R. E., Eriksen, C. C., Fratantoni, D. M., & Perry, M. J. (2004). Underwater gliders for ocean research. *Marine Technology Society Journal*, *38*(2), 73–84.

Rudnick, D. L., Davis, R. E., & Sherman, J. T. (2016). Spray underwater glider operations. *Journal of Atmospheric and Oceanic Technology*, *33*(6), 1113–1122.

Rudnick, D. L., Gopalakrishnan, G., & Cornuelle, B. D. (2015). Cyclonic eddies in the gulf of mexico: Observations by underwater gliders and simulations by numerical model. *Journal of Physical Oceanography*, *45*(1), 313–326.

Rudnick, D. L., Johnston, T. S., & Sherman, J. T. (2013). High-frequency internal waves near the luzon strait observed by underwater gliders. *Journal of Geophysical Research: Oceans*, *118*(2), 774–784.

Schönau, M. C., & Rudnick, D. L. (2015). Glider observations of the North Equatorial Current in the western tropical Pacific. *Journal of Geophysical Research: Oceans*, *120*(5), 3586–3605.

Schönau, M. C., & Rudnick, D. L. (2017). Mindanao current and undercurrent: Thermohaline structure and transport from repeat glider observations. *Journal of Physical Oceanography*, *47*(8), 2055–2075.

Sherman, J., Davis, R. E., Owens, W., & Valdes, J. (2002). The autonomous underwater glider" spray". *IEEE Journal of Oceanic Engineering*, *26*(4), 437–446.

Szekely, T., Gourrion, J., Pouliquen, S., Reverdin, G., & Merceur, F. (2019). *CORA, coriolis ocean dataset for reanalysis*.

Team, T. I. (2018). *International quality controlled ocean database (IQuOD) version 0.1 - aggregated and community quality controlled ocean profile data 1772-2018 (NCEI*

accession 0170893) [Data set]. NOAA National Centers for Environmental Information. https://doi.org/https://doi.org/10.7289/v51r6nsf

Todd, R. E., Owens, W. B., & Rudnick, D. L. (2016). Potential vorticity structure in the north atlantic western boundary current from underwater glider observations. *Journal of Physical Oceanography*, *46*(1), 327–348.

Tucker, T., Giglio, D., Scanderbeg, M., & Shen, S. S. (2020). Argovis: A web application for fast delivery, visualization, and analysis of argo data. *Journal of Atmospheric and Oceanic Technology*, *37*(3), 401–416.

Wong, A. P., Wijffels, S. E., Riser, S. C., Pouliquen, S., Hosoda, S., Roemmich, D., Gilson, J., Johnson, G. C., Martini, K., Murphy, D. J., & others. (2020). Argo data 1999–2019: Two million temperature-salinity profiles and subsurface velocity observations from a global array of profiling floats. *Frontiers in Marine Science*, *7*, 700.

Zeiden, K. L., Rudnick, D. L., & MacKinnon, J. A. (2019). Glider observations of a mesoscale oceanic island wake. *Journal of Physical Oceanography*, *49*(9), 2217–2235.