

CrocoLakeTools: A Python package to convert ocean observations to the parquet format

Enrico Milanese¹, David Nicholson¹, Gaël Forget², and Susan Wijffels¹

¹ Woods Hole Oceanographic Institution, United States of America ² Massachusetts Institute of Technology, United States of America

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [Open Journals](#)

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Investigations of the ocean state are possible thanks to the ever growing number of measurements performed with multiple instruments by different research missions. The vast and heterogeneous efforts have brought the community to define data storage conventions (e.g. CF-netCDF) and to assemble collections of datasets (e.g. the World Ocean Database). Yet, accessing these datasets often requires the usage of specialized tools, is generally inefficient over the cloud, and remains a major obstacle to new users in terms of pre-required knowledge and time. CrocoLakeTools is a Python package that addresses those challenges by providing workflows to convert various oceanographic datasets from their original format to a uniform parquet dataset with a shared schema.

Statement of need

CrocoLakeTools is a Python package to build workflows that convert ocean observations from their original formats (e.g. netCDF, CSV) to parquet. CrocoLakeTools takes advantage of Python's well-established and growing ecosystem of open tools: it uses dask's parallel computing capabilities to convert multiple files at once and to handle larger-than-memory data. dask is already well-integrated with xarray and pandas, two widely used Python libraries for the treatment of array and tabular data, respectively, and with pyarrow, the API to the Apache Arrow library which is used to generate the parquet dataset.

Parquet is a data storage format for big tidy data which presents several advantages: it is language agnostic (it can be accessed with Python, Matlab, Julia and web development technologies); it offers faster reading performances than other tabular formats such as CSV; it is optimized for cloud systems storage and operations; it is widespread in the data science community, leading to a multitude of freely accessible tools and educational material.

CrocoLakeTools was developed with the goal of building and serving CrocoLake, a regularly refreshed dataset containing sparse in-situ oceanographic observations that are pre-filtered to contain only quality-controlled measurements. CrocoLakeTools was designed to be used by researchers, engineers and data scientists in oceanography and to be accessed by the wider oceanographic community.

State of the field

In-situ measurements of ocean characteristics are collected by a wide range of projects at different scales: from small research groups to international consortia, from time-limited regional explorations to continuous global observation arrays. This has led to a highly heterogeneous landscape of products that often differ for schema (variable names, data types), data format,

and access services. As a result, ingesting and analyzing data from different sources, e.g. for research purposes, presents technical and logistical challenges. The need for homogenized ocean data products has then prompted several efforts, in particular when it comes to sparse observations, which are characterized by discrete measurements that are generally irregular and sparse in time and space.

One of the most comprehensive projects to date is the World Ocean Database (WOD) (Mishonov et al., 2024), which contains more than 40 variables gathered from several sources (buoys, gliders, floats, bathythermographs, etc.). The earliest record dates back to 1772, the data is quality-controlled, and the database is regularly updated with the latest measurements available. The data is public and freely available through different interfaces (WODselect, THREDDS, HTTPS, FTP) in ASCII, Comma Separated Value (CSV) and netCDF formats. WOD is maintained by the National Centers for Environmental Information (NCEI) of the United States' National Oceanic and Atmospheric Administration (NOAA).

Copernicus Marine Service is an initiative of the European Union that provides a wide range of ocean data products: physical, biogeochemical and sea ice data at regional and global scales, obtained from numerical models, in-situ observations and satellite observations. Among Copernicus' products, CORA (Szekely et al., 2019) (Coriolis Ocean database for ReAnalysis) provides quality-controlled in-situ temperature and salinity measurements gathered from several global and regional networks.

The International Quality-controlled Ocean Database (IQuOD) (Team, 2018) is an effort by the oceanographic community "with the goal of producing the highest quality and complete single ocean profile repository along with (intelligent) metadata and assigned uncertainties". It includes subsurface ocean profiles of several variables, paying particular attention to temperature measurements. The database is prepared by NCEI, and it is freely available in netCDF format through multiple channels (THREDDS, HTTPS, FTP).

The Ocean Data Platform by the non-profit HUB Ocean is among the youngest projects in the community, and it allows to find and access datasets from a catalog. The user can interact with the platform through different interfaces (SDK, REST API, OQS, JupyterHub workspaces), loading the datasets in tabular format.

Of particular interests to the community has been the data collected by the Argo (Wong et al., 2020) international program with a growing fleet of robotic floats, and several tools have been developed to handle such data. `argopy` (Maze & Balem, 2020) is a Python library that facilitates access and manipulation of data from the real-time observing system Argo (Wong et al., 2020). The user provides some filters (e.g. coordinate and time ranges, measurement name, etc) and `argopy` retrieves the relevant data from the Argo Global Data Assembly Center (GDAC). `ArgoData.jl` (Gael Forget, 2025) and `OceanRobots.jl` (Gael Forget & Gebbie, 2025) are Julia packages that provide similar functionalities to `argopy`. `Argovis` (Tucker et al., 2020) is a REST API and web application hosted at the University of Colorado, Boulder (United States) and serving profile data in JSON format from the Argo program, and from ship and drifters missions. `oce` (Kelley et al., 2022) is a package written in R providing multiple functions to access oceanographic observations stored in a variety of formats. It focuses on datasets for physical oceanography with the goal of providing more tools in the future. `argoFloats` (Kelley et al., 2021) is a derived package that focuses on fetching and manipulating Argo data.

The above efforts often serve data in ASCII, CSV, JSON, or netCDF formats. For context, netCDF is a binary format that offers the advantage to be compact and efficient when dealing with multidimensional data, while the others have the advantage of being human-readable (but can be very inefficient for large datasets). None of these formats is optimized for cloud object storage, although there are ongoing efforts for netCDF (e.g. Zarr and Icechunk). For this reason, Parquet has been drawing more and more attention from the earth sciences community recently.

Parquet is a cloud-optimized binary format for tidy and large data (i.e. large tables) that is

language-agnostic. It is widely used in the data science and corporate worlds, and the software ecosystem around it is sound, mature, and growing. Table 1 presents an overview of the characteristics of Parquet and the other formats. We chose Parquet as the target format for CrocoLake because (1) it is optimized for cloud storage and cloud computing, (2) its mature software ecosystem includes packages in multiple coding languages to access it (Python, Julia, MATLAB, web technologies, etc.), (3) its tabular format suits in-situ observations which are sparse and do not benefit from advantages of array-based formats as much as gridded datasets, and (4) novel users are generally more familiar with tabular data than with specialized oceanographic data formats.

As we aim to make CrocoLake easily accessible from a technical standpoint, the main drawbacks of Parquet at present are that (1) many workflows in ocean modeling are based on multidimensional data structures (not tabular ones) and (2) attaching attributes to the data is not straightforward (compared to netCDF). However, we see CrocoLake as a major building bloc in workflows that require fast access to sparse in-situ ocean observations, responding to necessities of data storage with fast access both on disk and the cloud, rather than as an end-all solution for ocean observations. For example, array-based storage formats might instead be more relevant to workflows that rely mostly on regular and gridded observations (such as satellite recordings).

Table 1. Comparison between different common file formats for oceanographic datasets.

Feature Name	CSV	netCDF	Parquet	ASCII	JSON	Zarr + Icechunk
Cloud-Optimized Structure	No Tabular	No Array-based	Yes Tabular	No Tabular	No Hierarchical	Yes Array-based
Available tools in:						
Python	Yes	Yes	Yes	Yes	Yes	Yes
Julia	Yes	Yes	Yes	Yes	Yes	Zarr only
MATLAB	Yes	Yes	Yes	Yes	Yes	Zarr only
Attributes descriptors	Dedicated columns, header	Dictionary, accessed with data	Dictionary, separate access from data	Dedicated column	Dedicated field in object	Dictionary, accessed with data

Another key feature of CrocoLakeTools is that it is fully open-source. Anyone can thus use CrocoLakeTools to build their own flavor of CrocoLake. We also welcome new contributors to CrocoLakeTools, for example by adding converters to support new datasets.

Software Design

Converters

The core task of CrocoLakeTools is to take one or more files from a dataset and convert them to parquet, ensuring that CrocoLake's schema is followed. This is achieved through the methods contained in the Converter class and its subclasses. While the conversion of all datasets requires some general functionality (e.g. renaming the original variables to the final schema), each conversion requires specialized tools that differ between datasets (e.g. the map used to rename variables). CrocoLakeTools then contains the Converter class, which contains the methods shared across datasets, and from which converter subclasses inherit and implement the specific needs of each dataset.

122 Workflow

123 Local mirrors

124 The first step in our workflow is to retrieve original files from each data provider (Figure 1).
 125 The original source data follow the format, schema, nomenclature, and conventions defined
 126 by their provider (project, mission, scientist, etc.) independently of CrocoLake's workflow.
 127 Modules to download the original data are optional. They are expected to inherit from the
 128 Downloader class and be called `downloader<DatasetName>` (e.g. `downloaderArgoGDAC`). At
 129 the time of writing CrocoLakeTools is released with a downloader to build a local mirror of the
 130 Argo Global Data Assembly Center (GDAC), and we hope to support more data providers in
 131 the future. Whether a downloader module exists or the user downloads the data themselves,
 132 the original data is stored on disk and this is the starting point for the converter.

133 Parquet datasets

134 The second step is to convert the data to parquet, and finally merge the datasets into
 135 CrocoLake (Figure 2). The core of CrocoLakeTools are the modules in the Converter class
 136 and its subclasses. Each original dataset has its own subclass called `converter<DatasetName>`,
 137 e.g. `converterGLODAP`; further specifiers can be added as necessary (e.g. at this time there a
 138 few different converters for Argo data to prepare different datasets). The need for a dedicated
 139 converter for each project despite the usage of common data formats (e.g. netCDF, CSV)
 140 is due to differences in schema (e.g. variable names or units). Depending on the dataset,
 141 multiple converters can be applied. For example, to create CrocoLake, Argo data goes through
 142 two converters: 1. `converterArgoGDAC`, which converts the original Argo GDAC preserving
 143 most of its original conventions; 2. `converterArgoQC`, which takes the output of the previous
 144 step and applies some filtering based on Argo's QC flags and makes the data conforming to
 145 CrocoLake's schema. At this time CrocoLakeTools is released with converters for data from
 146 Argo (Wong et al., 2020), GLODAP (Global Ocean Data Analysis Project, Key et al. (2015),
 147 Lauvset et al. (2016), Olsen et al. (2016)) and Spray Data (Sherman et al. (2002), Rudnick
 148 et al. (2004), Rudnick & Cole (2011), Gopalakrishnan et al. (2013), Johnston et al. (2013),
 149 Rudnick et al. (2013), Rudnick et al. (2015), Schönau & Rudnick (2015), Rudnick (2016),
 150 Rudnick et al. (2016), Todd et al. (2016), Centurioni et al. (2017), Schönau & Rudnick
 151 (2017), Zeiden et al. (2019), Heiderich & Todd (2020)). The choice of glider data is limited
 152 to Spray Data because they are provided in delayed mode (i.e. with comprehensive and manual
 153 quality-control): we are interested in including data from the wider IOOS gliders network in
 154 the future, yet to our knowledge quality-control processes for delayed mode data are not yet
 155 standardized across all the network. We are also working towards adding new datasets, such
 156 as observations collected by Saildrones used by the Ocean Climate Stations Project (NOAA)
 157 and expendable bathythermograph data collected by the Oleander Project.

158 CrocoLake

159 CrocoLake is one parquet dataset that contains all converted datasets merged together. This can
 160 be achieved with the script `merge_crocolake.py`. The script first creates a directory containing
 161 symbolic links to each converted dataset. It then uses the submodule `CrocoLakeLoader` to
 162 seamlessly load all the converted datasets into memory as one dask dataframe with a uniform
 163 schema, merges them into CrocoLake, and stores it back to disk.

164 CrocoLake can be accessed with several programming languages with just a few lines of codes:
 165 the submodules `CrocoLake-Python`, `CrocoLake-Matlab`, and `CrocoLake-Julia` contain tools and
 166 examples in some languages.




Figure 1: CrocoLake's workflow: 'downloader's. 'CrocoLakeTools' is set up to host modules that are dedicated to download the desired datasets from the web. It currently supports the download only of Argo data (solid line subset), and other datasets require the user to download them manually (dashed border subsets). Argo, Spray Data and GLODAP (Global Ocean Data Analysis Project) are the different data providers supported at the time of submission.




Figure 2: CrocoLake's workflow: 'converter's. 'converter's read the data in their original format, transform it following CrocoLake's conventions, converts it to parquet, and stores it back to disk. Each dataset is converted to its own parquet version. Thanks to the submodule 'CrocoLakeLoader', multiple parquet datasets are merged into a uniform dataframe which is saved to disk as CrocoLake. For each dataset, a version containing only physical variables ('PHY') or also biogeochemical variables ('BGC') can be generated.

167 Schema and metadata

168 The nomenclature, units and data types are generally based on Argo's (<https://vocab.nerc.ac.uk/collection/>)
169 and are detailed in the online documentation. For CrocoLake's variables that are not present
170 in the Argo program, we provide new names maintaining consistency with Argo's style.

171 The parquet format has two different types of metadata: metadata related to storage
172 (e.g. schema, number of rows, etc.) and column attributes. We use the latter type to
173 store information about the units for each variable present in CrocoLake. At the moment we
174 are not including other metadata from the original sources (e.g. information about principal
175 investigator, contact details, comments, etc.), but we are exploring solutions to include some
176 of this in the future, based also on feedback from users.

177 Profile numbering

178 Ocean data is often accessed by profiles and we provide this functionality for CrocoLake too:
179 the user can retrieve the profiles through the `CYCLE_NUMBER` variable, which is unique and
180 progressive for each `PLATFORM_NUMBER` of each subdataset (`DB_NAME`). As each original product
181 uses its own conventions, the `CYCLE_NUMBER` of some datasets is generated ad hoc during their
182 conversion if no obvious match with `CYCLE_NUMBER` exists. The procedure for each dataset is
183 detailed in the online documentation.

184 Quality control

185 CrocoLake contains only quality-controlled (QC) measurements. We rely exclusively on QCs
186 performed by the data providers and at the time we do not perform any additional QC ourselves
187 (although this might change in the future). Each parameter `<PARAM>` has a corresponding
188 `<PARAM>_QC` flag that is generally set to 1 to indicate that the data is reliable. For Argo
189 measurements, the original QC value is preserved, and only measurements with QC values of
190 1, 2, 5, and 8 considered.

191 Measurement errors

192 Each parameter `<PARAM>` has a corresponding `<PARAM>_ERROR` that indicates a measurement's
193 error as provided in the original dataset. When no error is provided, `<PARAM>_ERROR` is set to
194 null.

195 Research impact statement

196 CrocoLakeTools is used to build CrocoLake, a fast and cloud-optimized dataset of
197 oceanographic observations, making diverse observations accessible faster and more efficiently
198 as we show in the following.

199 Benchmarking

200 We here present a few examples of the advantages and disadvantages in accessing the generated
201 parquet datasets. We compare access to Argo data in a few different scenarios between
202 CrocoLake and `argopy` (Maze & Balem, 2020), which is widely used in the community. For
203 the comparison, we use the parquet Argo data generated by `converterArgoGDAC` (i.e. without
204 applying any filtering based on quality-control flags), as it is the version closest to the data
205 served by `argopy`. The parquet dataset is stored on the cloud in a Amazon Web Services
206 Simple Storage Solution (AWS S3) bucket, and `argopy` is called to access data from the Argo
207 GDAC, so the data is always accessed over the internet. The parquet dataset is accessed with
208 Python for comparison with `argopy`; we use `dask` but other packages exist to access parquet
209 data (e.g. `duckDB`, `polars`) and the choice of package can affect the performance.

The goal of this benchmarking exercise is to provide a ballpark estimate of reading performances in a few different scenarios to showcase the potential of including oceanographic datasets in parquet format in scientific workflows; comparison with an existing tool as argopy is qualitative to provide a reference to help evaluate the performance of the parquet-based workflow. More generally, the reading performance of a given dataset depends on multiple factors beyond the dataset format itself (coding language, selection parameters, local vs cloud-based access, network connectivity, storage service, etc.).

Benchmarks were performed on a Lenovo ThinkPad P14s Gen 4 equipped with an AMD Ryzen 7 PRO 7840U processor (8 cores, 16 threads, up to 5.13 GHz) and 26 GiB DDR4 RAM. Network connectivity was provided via Wi-Fi: bandwidth measurements indicated an average download speed of 466.41 Mbit/s (10 runs, range: 363.37–555.70 Mbit/s) and an upload speed of 27.34 Mbit/s (10 runs, range: 26.98–27.48 Mbit/s); latency to google.com averaged 26.2 ms (10 measurements, range: 20.9–27.6 ms). All non-essential applications were closed during benchmarking, and measurements were collected immediately prior to the tests to ensure reproducibility.

The following Tables 2 and 3 summarize data loading times for PHY and BGC Argo datasets from AWS S3 (Parquet) versus the argopy Python library. Benchmarks include space and time filtering and loading by float WMO ID.

Table 2 – North West Atlantic Regional Subsets (2017)

Data type	Subset size	Filter details	Parquet time	argopy time
PHY	Large	Lat: [25,60], Lon: [-80,-30], Date: 2017-01-01 to 2018-01-01	44.8 ± 0.4 s	1266 ± 8.6s
PHY	Small	Lat: [55,60], Lon: [-50,-55], Date: 2017-01-01 to 2018-01-01	19.5 ± 1.5 s	63 ± 1.9s
BGC	Large	Lat: [25,60], Lon: [-80,-30], Date: 2017-01-01 to 2018-01-01	8.5 ± 0.7 s	954 ± 1.1s
BGC	Small	Lat: [55,60], Lon: [-50,-55], Date: 2017-01-01 to 2018-01-01	5.0 ± 2.3 s	108 ± 4.3s

Table 3 – Loading Full Float Records

Data type	Floats	Parquet time	argopy time
PHY	1 float	4.2 ± 1.3 s	3.0 ± 0.3 s
PHY	20 floats	17.4 ± 0.8 s	22.5 ± 1.1 s
BGC	1 float	3.4 ± 1.8 s	6.2 ± 0.6 s
BGC	20 floats	8.0 ± 2.8 s	140.0 ± 6.0 s

Notes:

- All timings are mean ± standard deviation measured over ten runs, except for the “Large” subsets in Table 2 which are measured over two runs.
- Large/small regional subsets (Table 2) differ by latitude/longitude bounding boxes (see above).
- The floats of the records in Table 3 were randomly picked and are identified by the following WMO IDs:
 - PHY, single float: 6902801;

- 240 – PHY, 20 floats: 4902120, 6902707, 3901629, 6901177, 4901812, 4901765, 4901216,
241 4902322, 6902801, 6902753, 6902805, 5904749, 4901827, 4902419, 4902397,
242 4901755, 6902810, 5904007, 4902112, 1901596;
- 243 – BGC, single float: 4902410;
- 244 – BGC, 20 floats: 6901754, 6901030, 4902409, 6902810, 6901750, 6902805, 6901182,
245 6902686, 5904989, 6901603, 4902390, 5904770, 1901217, 6901751, 1901208,
246 6901752, 6902812, 6901023, 6901758, 6901524.

247 Documentation and updates

248 The [documentation](#) describes the specifics of each dataset (e.g. what quality-control filter we
249 apply to each dataset, the procedure to generate the profile numbers, etc.), and is updated
250 every time a new feature is made available.

251 This document reflects the code at the time of publication. Any future updates and changes
252 to the code will be documented in the code main branch, which will contain the most recent
253 and best version available of CrocoLakeTools.

254 Citation

255 If you use CrocoLakeTools and/or CrocoLake, please do not limit yourself to citing this
256 manuscript but also remember to cite the datasets that you have used as indicated in the
257 documentation. For example, if your work relies on Argo measurements, acknowledge Argo
258 ([Wong et al., 2020](#)). This is important both for the maintainers of each product to track their
259 impact and to acknowledge their efforts that made your work possible.

260 AI usage disclosure

261 AI usage was limited to code drafting via Github Copilot chat (models GPT-3 and GPT-
262 4) and based on human-made design decisions. The drafts were human-reviewed before
263 implementation and tested as the rest of the code. No AI agent was used, i.e. no AI tool
264 edited the codebase, the documentation, or the manuscript directly.

265 Acknowledgements

266 This material is based upon work supported by the National Science Foundation under Grant
267 No. 2311382, “Collaborative Research: Frameworks: A community platform for accelerating
268 observationally-constrained regional oceanographic modeling.” G.F. acknowledges funding from
269 NASA award 1686358, NOAA award 035724-00001, and ARIA’s POLEMIX award.

270 We acknowledge the Argo, Spray Data, and GLODAP programs for their efforts in collecting
271 ocean observations and making them publicly available. The Argo data are collected and made
272 freely available by the International Argo Program and the national programs that contribute
273 to it (<https://argo.ucsd.edu>, <https://www.ocean-ops.org>), and is part of the Global Ocean
274 Observing System. E.M. thanks Robert Todd for guidance in accessing Spray Data.

275 References

- 276 Centurioni, L. R., Hormann, V., Talley, L. D., Arzeno, I., Beal, L., Caruso, M., Conry, P.,
277 Echols, R., Fernando, H. J., Giddings, S. N., & others. (2017). Northern arabian sea
278 circulation-autonomous research (NASCar) a research initiative based on autonomous
279 sensors. *Oceanography*, 30(2), 74–87. <https://doi.org/10.5670/oceanog.2017.224>

- 280 Forget, Gael. (2025). *Euroargodev/ArgoData.jl* [Data set]. Zenodo. <https://doi.org/10.5281/zenodo.3633717>
- 281
- 282 Forget, Gaël, & Gebbie, G. J. (2025). *JuliaOcean/OceanRobots.jl* [Data set]. Zenodo. <https://doi.org/10.5281/zenodo.4718472>
- 283
- 284 Gopalakrishnan, G., Cornuelle, B. D., Hoteit, I., Rudnick, D. L., & Owens, W. B. (2013). State estimates and forecasts of the loop current in the gulf of mexico using the MITgcm and its adjoint. *Journal of Geophysical Research: Oceans*, 118(7), 3292–3314. <https://doi.org/10.1002/jgrc.20239>
- 285
- 286
- 287
- 288 Heiderich, J., & Todd, R. E. (2020). Along-stream evolution of gulf stream volume transport. *Journal of Physical Oceanography*, 50(8), 2251–2270. <https://doi.org/10.1175/JPO-D-19-0303.1>
- 289
- 290
- 291 Johnston, T. S., Rudnick, D. L., Alford, M. H., Pickering, A., & Simmons, H. L. (2013). Internal tidal energy fluxes in the south china sea from density and velocity measurements by gliders. *Journal of Geophysical Research: Oceans*, 118(8), 3939–3949. <https://doi.org/10.1002/jgrc.20311>
- 292
- 293
- 294
- 295 Kelley, D. E., Harbin, J., & Richards, C. (2021). argoFloats: An r package for analyzing argo data. *Frontiers in Marine Science*, 8, 635922. <https://doi.org/10.3389/fmars.2021.635922>
- 296
- 297 Kelley, D. E., Richards, C., & Layton, C. (2022). Oce: An r package for oceanographic analysis. *Journal of Open Source Software*, 7(71), 3594. <https://doi.org/10.21105/joss.03594>
- 298
- 299 Key, R. M., Olsen, A., Heuven, S. van, Lauvset, S. K., Velo, A., Lin, X., Schirnick, C., Kozyr, A., Tanhua, T., Hoppema, M., & others. (2015). Global ocean data analysis project, version 2 (GLODAPv2). *Ornl/Cdiac-162, Ndp-093*.
- 300
- 301
- 302 Lauvset, S. K., Key, R. M., Olsen, A., Van Heuven, S., Velo, A., Lin, X., Schirnick, C., Kozyr, A., Tanhua, T., Hoppema, M., & others. (2016). A new global interior ocean mapped climatology: The 1× 1 GLODAP version 2. *Earth System Science Data*, 8(2), 325–340. <https://doi.org/10.5194/essd-8-325-2016>
- 303
- 304
- 305
- 306 Maze, G., & Balem, K. (2020). Argopy: A python library for argo ocean data analysis. *Journal of Open Source Software*, 5. <https://doi.org/10.21105/joss.02425>
- 307
- 308 Mishonov, A. V., Boyer, T. P., Baranova, O. K., Bouchard, C. N., Cross, S. L., Garcia, H. E., Locarnini, R. A., Paver, C. R., Reagan, J. R., Wang, Z., & others. (2024). *World ocean database 2023*. <https://doi.org/10.25923/z885-h264>
- 309
- 310
- 311 Olsen, A., Key, R. M., Van Heuven, S., Lauvset, S. K., Velo, A., Lin, X., Schirnick, C., Kozyr, A., Tanhua, T., Hoppema, M., & others. (2016). The global ocean data analysis project version 2 (GLODAPv2)—an internally consistent data product for the world ocean. *Earth System Science Data*, 8(2), 297–323. <https://doi.org/10.5194/essd-8-297-2016>
- 312
- 313
- 314
- 315 Rudnick, D. L. (2016). Ocean research enabled by underwater gliders. *Annual Review of Marine Science*, 8, 519–541. <https://doi.org/10.1146/annurev-marine-122414-033913>
- 316
- 317 Rudnick, D. L., & Cole, S. T. (2011). On sampling the ocean using underwater gliders. *Journal of Geophysical Research: Oceans*, 116(C8). <https://doi.org/10.1029/2010jc006849>
- 318
- 319 Rudnick, D. L., Davis, R. E., Eriksen, C. C., Fratantoni, D. M., & Perry, M. J. (2004). Underwater gliders for ocean research. *Marine Technology Society Journal*, 38(2), 73–84. <https://doi.org/10.4031/002533204787522703>
- 320
- 321
- 322 Rudnick, D. L., Davis, R. E., & Sherman, J. T. (2016). Spray underwater glider operations. *Journal of Atmospheric and Oceanic Technology*, 33(6), 1113–1122. <https://doi.org/10.1175/jtech-d-15-0252.1>
- 323
- 324
- 325 Rudnick, D. L., Gopalakrishnan, G., & Cornuelle, B. D. (2015). Cyclonic eddies in the gulf of mexico: Observations by underwater gliders and simulations by numerical model. *Journal*
- 326

Milanese et al. (2026). CrocoLakeTools: A Python package to convert ocean observations to the parquet format. *Journal of Open Source Software*, 11(1), 1-10. <https://doi.org/10.1093/jos/11.1.1>